

# Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

20.3.2018



Lecture is every week Wednesday 09:00 - 11:00.

06.03.2019: topic, teams

13.03.2019: TISS registration, initial PR

20.03.2019: other registrations, guest lecture

27.03.2019: PR for first issue done, second started,  
HS: kleiner Schiffbau

03.04.2019: first issue done, PR for second

10.04.2019: mid-term submission of exercises

08.05.2019: (HS?)

15.05.2019:

22.05.2019:

29.05.2019:

05.06.2019: final submission of exercises

12.06.2019:

19.06.2019: last corrections of exercises

26.06.2019: exam

# Popular Topics

14 tools	4 design
9 testability	4 cascading
9 code-generation	4 architecture of access
7 context-awareness	3 <b>configuration sources</b>
6 specification	3 config-less systems
6 misconfiguration	2 secure conf
6 <b>complexity reduction</b>	2 architectural decisions
5 validation	1 push vs. pull
5 points in time	1 infrastructure as code
5 error messages	1 full vs. partial
5 auto-detection	1 convention over conf
4 user interface	1 CI/CD
4 introspection	0 documentation

# Tasks for today

(until 20.03.2019 23:59)

## Task

Registration for talk, homework, teamwork and make sure to say which programming languages you know in STUDENTS.ini.

## Task

Write text in at least one issue.

## Tasks for next week

(until 27.03.2019 23:59, hint: to get help submit at least one day earlier)

### Task

Description of homework as pull request in private repo. (Inside a folder for you, use GitHub name.)

### Task

Description of teamwork (which application, which CM tool) as pull request in private repo. (Inside a folder for your team.)

### Task

Fix at least one issue and write some text in at least one other issue.

# Environment Variables

- 1 Environment Variables
  - Requirements
  - Conclusion
- 2 Complexity
  - Trend
  - Calculation
  - Usage
- 3 Configuration Specification
  - Why?
  - How?
  - Visibility
  - Calculate Default Values
- 4 Architectural Decisions

# Semantics of Command-line Arguments (cont.)

- passed by main for a new process via  
(int argc, char \*\* argv)
- visible from other processes (e.g., via ps aux)
- could be passed along to subprocesses but hardly done
- need to be parsed by process
- portability: differences in parsing
- cannot be changed from outside (requires restart, no IPC)

# Semantics

- are also per-process (`/proc/self/environ`)
- are not visible from other processes
- are automatically inherited by subprocesses
- need to be parsed by process (`[extern] char **environ`) but API is provided (`getenv`)
- cannot be changed from outside (requires restart or an additional IPC mechanism)



## Task

What is wrong with the code in the book?

# getenv

- is widely standardized, including SVr4, POSIX.1-2001, 4.3BSD, C89, C99 [1],
- is supported by many programming languages, and
- enforces key=value convention.

# Usage

- ① bypassing other configuration accesses (Q: 45 %)
- ② locating configuration files
- ③ debugging and testing (Q: 55 %, S: 1,152, i. e. 43 %)
- ④ sharing configuration settings across applications (Q: 53 %, S: 716, i. e. 47 %)
- ⑤ for configuration settings unlikely to be changed by a user (Q: 20 %)
- ⑥ *“even when it is used inside a loop”* (Q: 2 %)

# Portability


- no separators for values defined
- case sensitivity problems
- often many environment variables for the same purpose: TMP, TEMP, or TMPDIR
- sometimes one environment variable for different purposes: PATH

How can we deal with the many sources?

### Requirement

*A configuration library must support all three popular ways for configuration access: configuration files, command-line options, and environment variables.*

# Cascading



## Task

Discuss the differences of mounting and cascading with your neighbor.

# User View

- command-line for trying out configuration settings
- environment variables for configuration settings within a shell
- configuration files for persistent configuration settings




# Conclusion

- three different configuration sources widely used
- all three used for different reasons but often for the same configuration settings
- many different configuration file formats
- abstractions: key-value, mounting, and cascading


# Complexity

- 1 Environment Variables
  - Requirements
  - Conclusion
- 2 Complexity
  - Trend
  - Calculation
  - Usage
- 3 Configuration Specification
  - Why?
  - How?
  - Visibility
  - Calculate Default Values
- 4 Architectural Decisions





# Trend Firefox



# Trend Chromium



# Trend Configuration Files



Xu et al. [9]

# Types of Complexity

- complexity in access:
  - many different formats
  - non-uniformity
  - transformations
- configuration settings
  - number of settings  $s$
  - number of values  $n$
  - dependences between settings

# Calculation of Complexity

Using enumerative combinatorics:

- number of configurations:  $n^s$
- for  $N$  groups of different  $n$  and  $s$  (i.e.,  $n_1 \dots n_N$  with  $s_1 \dots s_N$  occurrences):

$$\prod_{i=1}^N n_i^{s_i}$$

- more difficult to calculate (or unbounded) for dependences, module instantiations, arrays, ...

# Calculation of Complexity

Examples:

- 600 boolean settings in Apache httpd (let us assume  $n = 2$ ):  
 $2^{600} \approx 10^{180}$
- 19 integer settings:  $2^{32 \cdot 19} = 2^{608} \approx 10^{183}$
- 2000 boolean settings in Firefox [5]:  $2^{2000} \approx 10^{602}$



# Calculation of Complexity (cont.)

Examples:

- for 20 boolean and 20 enums with 5 possibilities:

$$2^{20} * 5^{20} = 10^{20}$$

- MySQL has 461 settings, of which 216 are non-simple types [9]  
(let us assume  $n = \{3, 20\}$ ):  $3^{245} * 20^{216} \approx 10^{397}$   
(settings are explained in 5560 pages<sup>1</sup>)
- an array with 1 – 20 boolean settings:  $2^{20}$

---

<sup>1</sup><https://downloads.mysql.com/docs/refman-5.7-en.pdf>

### Task

Calculate complexity for some tool you know.

### Task

Possible Homework: Write tool to calculate complexity with a given configuration specification.

# Decision Tree

- configuration settings may depend on each other
- form a decision tree [2, 8]
- the decision tree is an instantiation of chosen configuration settings
- calculation only needs to consider instantiations which make a difference:  
essential configuration complexity [6]

# Harmful Defaults [9]

- Problem: Two major data losses on a dozen machines.
- Cause: Stayed with the default values of the data-path settings (e.g., `dfs.name.dir`, `dfs.data.dir`) which point to locations in `/tmp`. Thus, after the machines reboot, data losses occur. “One of the common problems from users.” (from Cloudera)
- up to 53 % of misconfigurations is due to staying at defaults
- 17 % to 48 % of configuration issues are about difficulties in finding settings

# Unnecessary Settings [9]

- Configuration Parameter:  
`dfs.namenode.tolerate.heartbeat.multiplier`
- Developers' Discussion: Since we are not sure what is a good choice, how about making it configurable? We should add a configuration option for it. Even if it's unlikely to change, if someone does want to change it they'll thank us that they don't have to change the code/recompile to do so.
- Real-World Usage:
  - No usage found by searching the entire mailing lists and Google.
  - No usage reported in a survey of 15 Hadoop users in UCSD.

# Unnecessary Settings [9]

- 6 % to 17 % of settings set by majority
- up to 54 % are seldom set
- up to 47 % of numeric settings have no more than five distinct values

# Reduction

Q: *"Why do you think configuration should be reduced?"*

- to simplify code maintenance (50 %),
- to prevent errors and misconfiguration (43 %),
- to provide better user experience (40 %),
- ***"I do not think it should be reduced"*** (30 %),
- because they prefer auto-detection (29 %)  
(with a possibility to override configuration settings: 32 %),
- *"because use-cases which are rarely used should not be supported"* (13 %),
- *"never find time for this task"* (9 %), and
- *"because only standard use-cases should be supported"* (1 %)

### Question

How to specify reduction strategies of configuration settings?

### Answer

Configuration Specification



# Configuration Specification

- 1 Environment Variables
  - Requirements
  - Conclusion
- 2 Complexity
  - Trend
  - Calculation
  - Usage
- 3 Configuration Specification
  - Why?
  - How?
  - Visibility
  - Calculate Default Values
- 4 Architectural Decisions

# Rationale

- without specification you and others do not even know which settings are available
- needed for any further techniques we will discuss
- essential for *no-futz computing* Holland et al. [4]
- the foundation for any advanced tooling like configuration management tools
- needed as communication of producers and consumers of configuration

## Task

Brainstorming: What can be part of a configuration specification?

## Task

Advantages/Disadvantages?

## Task

Alternatives?

*Q: "Configuration specification (e.g. XSD/JSON schemas) allows you to describe possible values and their meaning. Why do/would you specify configuration?"*

- 58 % for *"looking up what the value does"*,
- 51 % it helps users to avoid common errors (*"so that users avoid common errors"*),
- 46 % to simplify maintenance,
- 40 % for rigorous validation,
- 39 % for documentation generation (for example, man pages, user guide),
- 30 % for external tools accessing configuration,
- 28 % for generating user interfaces,
- 25 % for code generation, and
- 24 % for specification of links between configuration settings.

Why?

# Limitations of Schemata designed for Data

- like XSD/JSON schemas
- they are already very helpful but:
- not key-value based
- not easy to introspect
- designed to validate data without semantics:  
file path vs. presence of file
- not always possible to extend with plugins
- tied to specific formats (e.g. XML/JSON)

Why?

# Limitations of Zero-Configuration

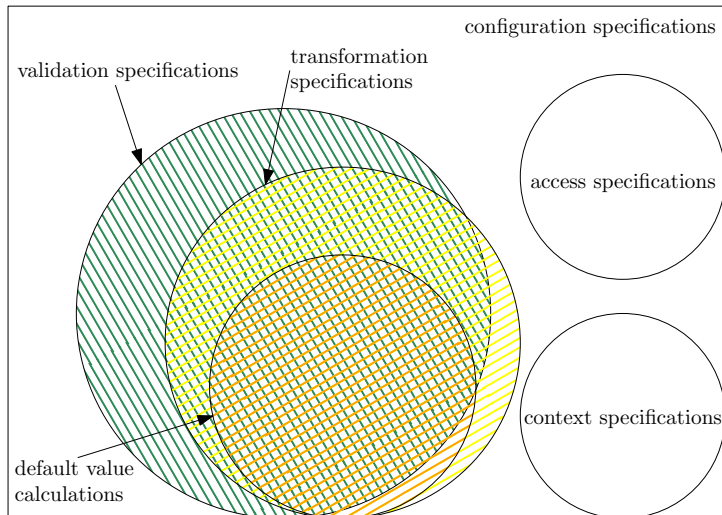
- e.g. gpsd<sup>1</sup>
- broken hardware or protocols
- auto-detection may go wrong
- the configuration actually lives elsewhere (e.g., in the GPS devices)

---


<sup>1</sup>[www.aosabook.org/en/gpsd.html](http://www.aosabook.org/en/gpsd.html)

How?

# Types of Specifications



# Metalevels





### Task

What do we mean with a configuration specification?

### Task

Which requirements do we have for a configuration specification?

# Requirements

- formal/informal?
- complete?
- should be extensible
- should be external to application
- open for introspection (for tooling)
- should talk to users
- should allow generation of artefacts

# Grammar

$\langle \text{configuration specifications} \rangle ::= \{ \langle \text{configuration specification} \rangle \}$

$\langle \text{configuration specification} \rangle ::= '[' \langle \text{key} \rangle ']' \langle \text{properties} \rangle$

$\langle \text{properties} \rangle ::= \{ \langle \text{property} \rangle \}$

$\langle \text{property} \rangle ::= \langle \text{property name} \rangle ':' [ \langle \text{property value} \rangle ]$

How?

# Example

```
1 [slapd/threads/listener]
2 default := 1
3 type := int
```

# Visibility

- idea: show only relevant settings for specific user group
- or disallow editing: accessibility
- requires user-feedback loops [9]
- most-used settings should be best visible (or even enforce them to be changed: against harmful defaults)
- think of your users (administrators),  
only expose what users need
- write an rationale why someone needs it

# Example

```
1 [slapd/threads/listener]
2 visibility := developer
3
4 [slapd/access/#]
5 visibility := user
```

## Task

Brainstorming: Now, how do we implement such a specification?

# Implementations

For example:

- generate examples/documentation
- auto-completion/syntax highlighting/IDE support
- tooling (GUI, Web UI)
- validate configuration files
- visudo-like
- plugins in configuration framework



- idea: make default value better
- is the generalization of sharing configuration values
- can be combined with visibility
- can be derived from other configuration settings
- can be derived from context [7]
- can be derived from hardware/system (problem with dependences)
- XServer vs. gpsd

# Examples

Sharing:

```
1 [slapd/threads/listener]
2 fallback/#0 :=slapd/threads
```

Percentages

(e.g., configured image should be additionally cropped):

```
1 [image/width]
2 type :=int
3
4 [crop]
5 type :=int
6 check/range :=0-100
```

# Examples

Context:

```
1 [slapd/threads/listener]
2 context := /slapd/threads/%cpu%/listener
```

Calculation with Context

(e.g., switch off GPS if battery low):

```
1 [gps/status]
2 assign := (battery > 'low') ? ('on') : ('off')
```

## Architectural Decisions

- 1 Environment Variables
  - Requirements
  - Conclusion
- 2 Complexity
  - Trend
  - Calculation
  - Usage
- 3 Configuration Specification
  - Why?
  - How?
  - Visibility
  - Calculate Default Values
- 4 Architectural Decisions

# Software Architecture

- architecture is high-level description of the overall system
- use ready-made patterns and templates for architecture
- e.g., <http://arc42.org/>
- architectural decisions [3] essential (e.g., Chapter 9 in arc42)

# Architectural Decisions

- describe decisions that lead to the architecture
- open decisions are high-level configuration
- useful to have patterns [3] and templates, too
- template: problem, constraints, assumptions, considered alternatives, decision, rationale, implications, related, notes

## Why are configuration settings added?

The typical reasons are:

- 1 a requirement,
- 2 an architectural decision,
- 3 a technical need, and
- 4 an ad hoc decision.

# in Configuration Specification

```
1 [slapd/threads/listener]
2 description := adjust to use more threads
3 rationale := needed for many-core systems
4 requirement := 1234
5 visibility := developer
```



# Conclusion

- alarming trend in number and complexity of configuration settings
- sharing, visibility and default value calculation often helps
- needs abstraction: configuration specification
- but also more courageous decisions and periodical reevaluation
- different ways to reduce configuration space

- [1] *getenv(3) Linux User's Manual*, March 2017.
- [2] Krzysztof Czarnecki, Paul Grünbacher, Rick Rabiser, Klaus Schmid, and Andrzej Wąsowski. Cool features and tough decisions: A comparison of variability modeling approaches. In *Proceedings of the Sixth International Workshop on Variability Modeling of Software-Intensive Systems, VaMoS '12*, pages 173–182, New York, NY, USA, 2012. ACM. ISBN 978-1-4503-1058-1. doi: 10.1145/2110147.2110167. URL <http://dx.doi.org/10.1145/2110147.2110167>.
- [3] Neil B Harrison, Paris Avgeriou, and Uwe Zdun. Using patterns to capture architectural decisions. *Software, IEEE*, 24(4):38–45, 2007. ISSN 0740-7459. doi: 10.1109/MS.2007.124.

- [4] David A. Holland, William Josephson, Kostas Magoutis, Margo I. Seltzer, Christopher A. Stein, and Ada Lim. Research issues in no-futz computing. In *Hot Topics in Operating Systems, 2001. Proceedings of the Eighth Workshop on*, pages 106–110. IEEE, May 2001. doi: 10.1109/HOTOS.2001.990069.
- [5] Dongpu Jin, Xiao Qu, Myra B. Cohen, and Brian Robinson. Configurations everywhere: Implications for testing and debugging in practice. In *Companion Proceedings of the 36th International Conference on Software Engineering, ICSE Companion 2014*, pages 215–224, New York, NY, USA, 2014. ACM. ISBN 978-1-4503-2768-8. doi: 10.1145/2591062.2591191. URL <http://dx.doi.org/10.1145/2591062.2591191>.

- [6] J. Meinicke, C. P. Wong, C. Kästner, T. Thüm, and G. Saake. On essential configuration complexity: Measuring interactions in highly-configurable systems. In *2016 31st IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 483–494, Sept 2016.
- [7] Markus Raab and Gergö Barany. Introducing context awareness in unmodified, context-unaware software. In *Proceedings of the 12th International Conference on Evaluation of Novel Approaches to Software Engineering - Volume 1: ENASE,,* pages 218–225. INSTICC, ScitePress, 2017. ISBN 978-989-758-250-9. doi: 10.5220/0006326602180225.
- [8] Mark-Oliver Reiser. *Core Concepts of the Compositional Variability Management Framework (CVM): A Practitioner's Guide*. TU, Professoren der Fak. IV, 2009.

- [9] Tianyin Xu, Long Jin, Xuepeng Fan, Yuanyuan Zhou, Shankar Pasupathy, and Rukma Talwadker. Hey, you have given me too many knobs! Understanding and dealing with over-designed configuration in system software. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering, ESEC/FSE 2015*, pages 307–319, New York, NY, USA, 2015. ACM. ISBN 978-1-4503-3675-8. doi: 10.1145/2786805.2786852. URL <http://dx.doi.org/10.1145/2786805.2786852>.