


# Configuration Management

Markus Raab

Institute of Information Systems Engineering, TU Wien

27.4.2018



# Organization

Next dates:

27.4.2018: **concepts for team exercise**

4.5.2018: lecture

18.5.2018: guest lecture

25.5.2018: team exercise submitted

1.6.2018: lecture

8.6.2018: lecture

15.6.2018: last corrections of team exercise

22.6.2018: test

## Popular Topics

4 validation	2 configuration specification
4 user interface	2 command-line args
3 tools (benefits?)	2 code generation
3 testability	1 variability
3 complexity reduction (when conf. needed?)	1 self-description
3 architectural decisions	1 round-tripping
2 Puppet	1 early
2 modularity	1 introspection
2 environment variables	1 dependences
2 documentation	1 auto-detection
	1 context-awareness
	1 administrators

# Configuration Access (Recapitulation)

## Configuration Access (Recapitulation)

**Configuration access** is the part of every software system concerned with fetching and storing configuration settings from and to the execution environment. There are many ways to access configuration [3, 5, 9]. **Configuration access APIs** are APIs that enable configuration access.

Within the source code the **configuration access points** are configuration access API invocations that return configuration values.

## Trend (Recapitulation)

## Trend (Recapitulation)

- alarming trend in number and complexity of configuration settings
- sharing, visibility and default value calculation often helps
- needs abstraction: configuration specification
- but also more courageous decisions and periodical reevaluation
- different ways to reduce configuration space

# SpecElektra (Recapitulation)



## SpecElektra (Recapitulation)

SpecElektra is a modular ***configuration specification language*** for configuration settings. In SpecElektra we use properties to specify configuration settings and configuration access. SpecElektra enables us to specify different parts of Elektra.


# Modularity (Recapitulation)

## Modularity (Recapitulation)

*Vertical modularity* describes how strongly separated the configuration accesses of different applications is.

*Horizontal modularity* describes how strongly separated modules implementing configuration access for a single application is.

# Vertical Modularity (Recapitulation)



Needed to keep applications independently. Boxes are applications, cylinders are configuration files, F? are frontends or frontend adapters, L? are configuration libraries [7].

## Plugins (Recapitulation)

## Plugins (Recapitulation)

**Plugins** are filters, sinks, and sources processing a key set. We aim at SpecElektra to be as modular as possible and make extensive use of plugins:

- 1 SpecElektra does not have any built-in feature, all features are (or can be) implemented as plugins.
- 2 Elektra works completely without SpecElektra's specifications.
- 3 Configuration specifications are present within the execution environment. Thus any tool and plugin can introspect and use the specifications.

# Conclusion

- Challenges: duplications, transformations, ...
- KeySet equivalence: settings are instantiated configuration files
- Configuration access APIs with code generation
- Guarantees of configuration access points
- We reuse properties of SpecElektra (type, default)
- We prefer hierarchies and tags to long function names
- Usually introspection preferred, except for static type safety

## Code Generation vs. Introspection

- 1 Code Generation vs. Introspection
- 2 Testability
- 3 Early Detection



# Introspection (Recapitulation)

# Introspection (Recapitulation)

- unified get/set access to (meta\*)-key/values
- access via applications, CLI, GUI, web-UI, ...
- access via any programming language (similar to file systems)
- GUI, web-UI can semantically interpret metadata

# Rationale (Recapitulation)

## Task

How to ensure that configuration access points match with present configuration settings?

# Rationale (Recapitulation)

## Task

How to ensure that configuration access points match with present configuration settings?

### Configuration Specification:

- without specification you and others do not even know which settings are available
- needed for any further techniques we will discuss:
  - code generation guarantees that configuration access points match with specification
  - validation guarantees that configuration settings match with specification

# Internal Specification

For example, OWNER:

```
1 import org.aeonbits.owner.Config;
2
3 public interface ServerConfig extends Config {
4     int port();
5     String hostname();
6     @DefaultValue("42")
7     int maxThreads();
8 }
```

## Task

Why do we need an external specification?

## Task

Why do we need an external specification?

### Introspection:

- essential for *no-futz computing* Holland et al. [2]
- the foundation for any advanced tooling like configuration management tools
- needed as communication of producers and consumers of configuration

# External Specification

```
1 [port]
2 type := long
3 [hostname]
4 default := 42
5 [threads/max]
6 type := long
```

- read and writable by other applications
- we still can generate the internal specification
- furthermore, we fulfill needs for configuration management tools



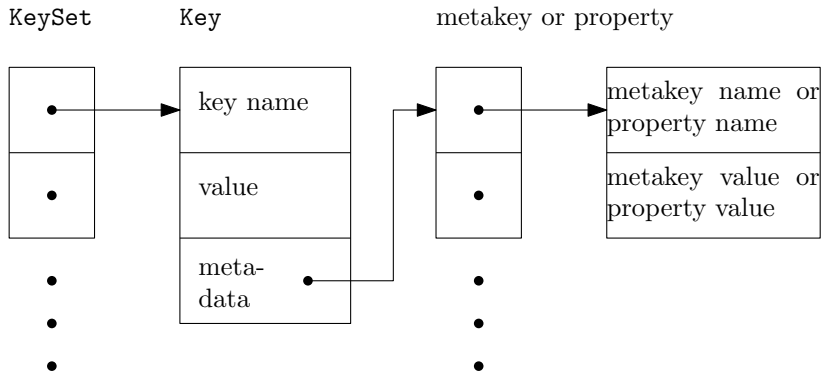
Other Artefacts (Recapitulation):

## Other Artefacts (Recapitulation):

- examples (e.g., defaults)
- documentation
- auto-completion/syntax highlighting/IDE support
- tooling (GUI, Web UI)
- validation code
- configuration management tool code

# KeySet (Recapitulation)

The common data structure between plugins:



# KeySet Generation (Recapitulation)

## Question

Idea: What if the configuration file format grammar describes source code?

# KeySet Generation (Recapitulation)

## Question

Idea: What if the configuration file format grammar describes source code?

key spec: /slapd/threads/listener, with the configuration value 4 and the property default  $\mapsto$  1:

```
1 ksNew (keyNew ("spec:/slapd/threads/listener",
2               KEY_VALUE, "4",
3               KEY_META, "default", "1",
4               KEY_END),
5       KS_END);
```

## Finding

We get source code representing the settings.

# Possible Properties (Recapitulation)

## Possible Properties (Recapitulation)

For example, SpecElektra has following properties:

**type** represents the type to be used in the emitted source code.

**opt** is used for short command-line options to be copied to the namespace proc.

**opt/long** is used for long command-line options, which differ from short command-line options by supporting strings and not only characters.

**readonly** yields compilation errors when developers assign a value to a contextual value within the program.

**default** enables us to start the application even if the backend does not work.

# (Recapitulation)

## Question

Introspection vs. Code Generation?



# (Recapitulation)

## Question

Introspection vs. Code Generation?

## (Recapitulation)

### Question

Introspection vs. Code Generation?

- more techniques for performance improvements with code generation
- + specification can be updated live on the system without recompilation
- + tooling has generic access to all specifications
- + new features the key database (e.g., better validation) are immediately available consistently

### Implication

We generally prefer introspection, except for a very thin configuration access API.

# Testability

- 1 Code Generation vs. Introspection
- 2 Testability
- 3 Early Detection

## Question

What do we want to test?

## Question

What do we want to test?

- That settings do what they should (devs and admins)
- That settings are properly validated (devs [9])
- Regression tests [6]

## Question

What do we want to test?

- That settings do what they should (devs and admins)
- That settings are properly validated (devs [9])
- Regression tests [6]
- Are all settings implemented?
- Are all settings used in tests?
- Are there unused settings in the code?

Matt Welsh from Google wrote in 2013:<sup>1</sup>

*“Of course we have extensive testing infrastructure, but the ‘hard’ problems always come up when running in a real production environment, with real traffic and real resource constraints. Even integration tests and canarying are a joke compared to how complex production-scale systems are.”*

Most of these problems are still not well understood.

---

<sup>1</sup>What I wish systems researchers would work on. Retrieved from <http://matt-welsh.blogspot.com/2013/05/what-i-wish-systems-researchers-would.html>.

## Jin et al. [3]

- Wants to improve configuration-ware testing and debugging
- Manual investigations for three applications
- Finds 1957 settings in Firefox ( $2^{846} * 3^{1111}$ ) and 36322 in LibreOffice ( $2^{4433} * 3^{31889}$ )
- Finds unused settings: settings only in the source code
- Finds unsynchronized configuration settings (see in “early”)

## Requirement

*Configuration setting traceability is a necessity.*

## Idea

Code generation helps to trace settings and to find unused settings.



## Testing by developers:

- ConfErr [4] uses models of key board layout, psychology and linguistics. Tool injects possible misconfiguration.
- Spex Xu et al. [9] analyzes the source code to find misconfigurations. As by-product it extracts internal specifications (including transformation bugs).
- External specification can be directly used to generate test cases.

# Find Unused Settings

The first (optional) step of the algorithm is:

- Run all tests with code coverage.
- Check if generated code is executed.
- If it is, we know that the configuration setting is used in a test case. Otherwise, we know it is not tested by the test suite. All these untested configuration settings are remembered as candidates for the second step.

```
1 KeySet findUnusedSettings (KeySet untestedSettings,
2                             KDB kdb,
3                             Builder build)
4 {
5     KeySet unusedSettings = {};
6     KeySet configurationSpecification;
7     kdb.get (configurationSpecification);
8
9     for (candidate: untestedSettings)
10    {
11        configurationSpecification.remove (candidate);
12        kdb.set (configurationSpecification);
13        build.recompile ();
14        if (build.wasSuccessful ())
15        {
16            unusedSettings.append (candidate);
17        }
18        configurationSpecification.append (candidate);
19    }
20
21    kdb.set (configurationSpecification);
22    return unusedSettings;
23 }
```

## Early Detection

- 1 Code Generation vs. Introspection
- 2 Testability
- 3 Early Detection

# When are settings used?

- Implementation-time** configuration accesses are hard-coded settings in the source code repository. For example, architectural decisions [1] lead to implementation-time settings.
- Compile-time** configuration accesses are configuration accesses resolved by the build system while compiling the code.
- Deployment-time** configuration accesses are configuration accesses while the software is installed.
- Load-time** configuration accesses are configuration accesses during the start of applications.
- Run-time** configuration accesses are configuration accesses during execution not limited to the startup procedure.

# Latent Misconfiguration

Phases when we can detect misconfigurations:

- Compilation stage in configuration management tool
- Writing configuration settings on nodes
- Starting applications (load-time)
- When configuration setting is actually used (run-time)

## Problem

More context vs. easier to detect and fix.

As shown in [10]:

- 12 % – 39 % configuration settings are not used at all during initialization.

As shown in [10]:

- 12 % – 39 % configuration settings are not used at all during initialization.
- Applications often have latent misconfigurations (14 % – 93 %)



As shown in [10]:

- 12 % – 39 % configuration settings are not used at all during initialization.
- Applications often have latent misconfigurations (14 % – 93 %)
- Latent misconfigurations are particular severe (75 % of high-severity misconfigurations)

As shown in [10]:

- 12 % – 39 % configuration settings are not used at all during initialization.
- Applications often have latent misconfigurations (14 % – 93 %)
- Latent misconfigurations are particular severe (75 % of high-severity misconfigurations)
- Latent misconfiguration need longer to diagnose

## Example [10]

Squid uses `diskd_program` but not before requests are served.  
Latent misconfiguration caused 7h downtime and 48h diagnosis effort.

## Example [10]

Squid uses `diskd_program` but not before requests are served.  
Latent misconfiguration caused 7h downtime and 48h diagnosis effort.

### Finding

Configuration from all external programs need to be checked, too.

## Using code generation

Code generation makes sure that only specified configuration settings are used.

# Checkers as plugins

Using checkers as plugins exclude whole classes of errors such as:

- Invalid file paths using the plugin “*path*”.
- Invalid IP addresses or host names using the plugins “*network*” or “*ipaddr*”.

Because the checks occur before the resources are actually used, the checks are subject to race conditions.

For example, a path that was present during the check, can have been removed when the application tries to access it.

In some situations facilities of the operating system help,<sup>1</sup> in others we have fundamental problems.<sup>2</sup>

---

<sup>1</sup>For example, we open the file during the check and pass `/proc/<pid>/fd/<fd>` to the application. This file cannot be unlinked, but unfortunately the file descriptor requires resources.

<sup>2</sup>For example, if the host we want to reach has gone offline after validation.

# Conclusion

- provide external specifications for other tooling and configuration management

# Conclusion

- provide external specifications for other tooling and configuration management
- use code generation to keep internal specifications consistent with external specifications



# Conclusion

- provide external specifications for other tooling and configuration management
- use code generation to keep internal specifications consistent with external specifications
- implement checkers as plugins

# Conclusion

- provide external specifications for other tooling and configuration management
- use code generation to keep internal specifications consistent with external specifications
- implement checkers as plugins
- execute checkers as early as possible

- [1] Neil B Harrison, Paris Avgeriou, and Uwe Zdun. Using patterns to capture architectural decisions. *Software, IEEE*, 24(4): 38–45, 2007. ISSN 0740-7459. doi: 10.1109/MS.2007.124.
- [2] David A. Holland, William Josephson, Kostas Magoutis, Margo I. Seltzer, Christopher A. Stein, and Ada Lim. Research issues in no-futz computing. In *Hot Topics in Operating Systems, 2001. Proceedings of the Eighth Workshop on*, pages 106–110. IEEE, May 2001. doi: 10.1109/HOTOS.2001.990069.
- [3] Dongpu Jin, Xiao Qu, Myra B. Cohen, and Brian Robinson. Configurations everywhere: Implications for testing and debugging in practice. In *Companion Proceedings of the 36th International Conference on Software Engineering, ICSE Companion 2014*, pages 215–224, New York, NY, USA, 2014. ACM. ISBN 978-1-4503-2768-8. doi:

10.1145/2591062.2591191. URL

<http://dx.doi.org/10.1145/2591062.2591191>.

- [4] Lorenzo Keller, Prasang Upadhyaya, and George Candea. Conferr: A tool for assessing resilience to human configuration errors. In *Dependable Systems and Networks With FTCS and DCC, 2008.*, pages 157–166. IEEE, 2008.
- [5] Emre Kiciman and Yi-Min Wang. Discovering correctness constraints for self-management of system configuration. In *International Conference on Autonomic Computing, 2004. Proceedings.*, pages 28–35. IEEE, May 2004. doi: 10.1109/ICAC.2004.1301344.

- [6] Xiao Qu, Myra B. Cohen, and Gregg Rothermel. Configuration-aware regression testing: An empirical study of sampling and prioritization. In *Proceedings of the 2008 International Symposium on Software Testing and Analysis*, ISSTA '08, pages 75–86, New York, NY, USA, 2008. ACM. ISBN 978-1-60558-050-0. doi: 10.1145/1390630.1390641. URL <http://doi.acm.org/10.1145/1390630.1390641>.
- [7] Markus Raab. Improving system integration using a modular configuration specification language. In *Companion Proceedings of the 15th International Conference on Modularity*, MODULARITY Companion 2016, pages 152–157, New York, NY, USA, 2016. ACM. ISBN 978-1-4503-4033-5. doi: 10.1145/2892664.2892691. URL <http://dx.doi.org/10.1145/2892664.2892691>.

- [8] Markus Raab and Gergő Barany. Introducing context awareness in unmodified, context-unaware software. In *Proceedings of the 12th International Conference on Evaluation of Novel Approaches to Software Engineering - Volume 1: ENASE*, pages 218–225. INSTICC, ScitePress, 2017. ISBN 978-989-758-250-9. doi: 10.5220/0006326602180225.
- [9] Tianyin Xu, Jiaqi Zhang, Peng Huang, Jing Zheng, Tianwei Sheng, Ding Yuan, Yuanyuan Zhou, and Shankar Pasupathy. Do not blame users for misconfigurations. In *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles*, pages 244–259. ACM, 2013.

- [10] Tianyin Xu, Xinxin Jin, Peng Huang, Yuanyuan Zhou, Shan Lu, Long Jin, and Shankar Pasupathy. Early Detection of Configuration Errors to Reduce Failure Damage. In *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI'16)*, Savannah, GA, USA, November 2016.